

# main

November 28, 2024

## CAB FARE PREDICTION

### 0.1 Import libraries and load dataset

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

[2]: traindf = pd.read_csv("D:\Study\code\project\Cab-Fare-Prediction-master\cab.
↪ csv", na_values={"pickup_datetime": "43"})
traindf.head()
```

```
[2]:   fare_amount      pickup_datetime  pickup_longitude  pickup_latitude  \
0         4.5  2009-06-15 17:26:21 UTC        -73.844311         40.721319
1        16.9  2010-01-05 16:52:16 UTC        -74.016048         40.711303
2         5.7  2011-08-18 00:35:00 UTC        -73.982738         40.761270
3         7.7  2012-04-21 04:30:42 UTC        -73.987130         40.733143
4         5.3  2010-03-09 07:51:00 UTC        -73.968095         40.768008

      dropoff_longitude  dropoff_latitude  passenger_count
0        -73.841610         40.712278             1.0
1        -73.979268         40.782004             1.0
2        -73.991242         40.750562             2.0
3        -73.991567         40.758092             1.0
4        -73.956655         40.783762             1.0
```

### 0.2 Data clearning

```
[3]: traindf.isna()

[3]:   fare_amount  pickup_datetime  pickup_longitude  pickup_latitude  \
0          False          False          False          False
1          False          False          False          False
2          False          False          False          False
3          False          False          False          False
4          False          False          False          False
```

...	...	...	...	...
16062	False	False	False	False
16063	False	False	False	False
16064	False	False	False	False
16065	False	False	False	False
16066	False	False	False	False

	dropoff_longitude	dropoff_latitude	passenger_count
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...	...	...	...
16062	False	False	False
16063	False	False	False
16064	False	False	False
16065	False	False	False
16066	False	False	True

[16067 rows x 7 columns]

```
[4]: traindf.describe()
```

```
[4]:
```

	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	\
count	16067.000000	16067.000000	16067.000000	16067.000000	
mean	-72.462787	39.914725	-72.462328	39.897906	
std	10.578384	6.826587	10.575062	6.187087	
min	-74.438233	-74.006893	-74.429332	-74.006377	
25%	-73.992156	40.734927	-73.991182	40.734651	
50%	-73.981698	40.752603	-73.980172	40.753567	
75%	-73.966838	40.767381	-73.963642	40.768014	
max	40.766125	401.083332	40.802437	41.366138	

	passenger_count
count	16012.000000
mean	2.625070
std	60.844122
min	0.000000
25%	1.000000
50%	1.000000
75%	2.000000
max	5345.000000

### 0.2.1 Changing the dtype

```
[5]: traindf['fare_amount'] = pd.to_numeric(traindf['fare_amount'], errors='coerce')
```

```
[6]: traindf['pickup_datetime'] = pd.to_datetime(traindf['pickup_datetime'],  
        ↪format='%Y-%m-%d %H:%M:%S UTC')
```

### 0.2.2 Extract Hour, Date, Day, Month, Year

```
[7]: traindf['Hour'] = traindf['pickup_datetime'].dt.hour  
traindf['minute'] = traindf['pickup_datetime'].dt.minute  
traindf['date'] = traindf['pickup_datetime'].dt.day  
traindf['day'] = traindf['pickup_datetime'].dt.dayofweek  
traindf['month'] = traindf['pickup_datetime'].dt.month  
traindf['year'] = traindf['pickup_datetime'].dt.year
```

```
[8]: traindf.head()
```

```
[8]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	\
0	4.5	2009-06-15 17:26:21	-73.844311	40.721319	
1	16.9	2010-01-05 16:52:16	-74.016048	40.711303	
2	5.7	2011-08-18 00:35:00	-73.982738	40.761270	
3	7.7	2012-04-21 04:30:42	-73.987130	40.733143	
4	5.3	2010-03-09 07:51:00	-73.968095	40.768008	

	dropoff_longitude	dropoff_latitude	passenger_count	Hour	minute	date	\
0	-73.841610	40.712278	1.0	17.0	26.0	15.0	
1	-73.979268	40.782004	1.0	16.0	52.0	5.0	
2	-73.991242	40.750562	2.0	0.0	35.0	18.0	
3	-73.991567	40.758092	1.0	4.0	30.0	21.0	
4	-73.956655	40.783762	1.0	7.0	51.0	9.0	

	day	month	year
0	0.0	6.0	2009.0
1	1.0	1.0	2010.0
2	3.0	8.0	2011.0
3	5.0	4.0	2012.0
4	1.0	3.0	2010.0

## 0.3 Handle missing values

### SimpleImputer

#### 0.3.1 drop null row from pickup\_datetime

```
[9]: traindf.drop(traindf[traindf['pickup_datetime'].isna()].index, inplace=True)
```

```
[10]: traindf.drop('pickup_datetime', axis=1, inplace=True)
```

```
[11]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
#apply it to features

imputer.fit(traindf[['fare_amount', 'passenger_count']])
traindf[['fare_amount', 'passenger_count']] = imputer.
    ↪transform(traindf[['fare_amount', 'passenger_count']])
```

```
[12]: traindf.isna().sum()
```


```
[12]: fare_amount          0
pickup_longitude         0
pickup_latitude          0
dropoff_longitude        0
dropoff_latitude         0
passenger_count          0
Hour                     0
minute                   0
date                     0
day                      0
month                    0
year                     0
dtype: int64
```

```
[13]: traindf.shape
```

```
[13]: (16066, 12)
```


```
[14]: plt.figure(figsize=(10,5))
sns.boxplot(y=traindf['passenger_count'])
```

```
[14]: <Axes: ylabel='passenger_count'>
```



```
[15]: plt.figure(figsize=(10,5))
      sns.boxplot(y=traindf['fare_amount'])
```

```
[15]: <Axes: ylabel='fare_amount'>
```



#### 0.4 Remove outlier

```
[16]: Q1 = traindf['passenger_count'].quantile(0.25)
      Q3 = traindf['passenger_count'].quantile(0.75)

      IQR = Q3 - Q1
```

```
lower = Q1 - 1.5*IQR
upper = Q3 + 1.5*IQR
print('Lower :',lower)
print('Upper :',upper)
```

```
Lower : -0.5
Upper : 3.5
```

### 1. Remove passenger\_count less than zero and with more than 6

```
[17]: traindf['passenger_count'].sort_values(ascending=True).head(10)
```

```
[17]: 13742    0.0
      2425    0.0
      6575    0.0
      5150    0.0
      4248    0.0
      5058    0.0
      9159    0.0
      3413    0.0
      4114    0.0
     15514    0.0
      Name: passenger_count, dtype: float64
```

```
[18]: traindf.drop(traindf[traindf['passenger_count'] > 4 ].index, inplace = True)
      traindf.drop(traindf[traindf['passenger_count'] < 0 ].index, inplace = True)
      traindf.drop(traindf[traindf['passenger_count'] == 0.12 ].index, inplace = True)
      traindf.drop(traindf[traindf['passenger_count'].isna() ].index, inplace = True)
```

### 2. Remove fare\_amount less than zero and should be less than 454

```
[19]: traindf['fare_amount'].sort_values(ascending=True).head(10)
```

```
[19]: 13032   -3.00
      2039   -2.90
      2486   -2.50
     10002    0.00
      2780    0.01
      1427    1.14
      8596    2.50
       503    2.50
      8711    2.50
      6002    2.50
      Name: fare_amount, dtype: float64
```

```
[20]: traindf.drop(traindf[traindf['fare_amount'] > 150 ].index, inplace = True)
      traindf.drop(traindf[traindf['fare_amount'] == 0 ].index, inplace = True)
      traindf.drop(traindf[traindf['fare_amount'].isna() ].index, inplace = True)
```

### 3. Pickup and dropoff latitude should be (-90 to 90 )

```
[21]: traindf.drop(traindf[traindf['pickup_latitude'] > 90].index, inplace=True)
      traindf.drop(traindf[traindf['pickup_latitude'] < -90].index, inplace=True)
      traindf.drop(traindf[traindf['dropoff_latitude'] > 90].index, inplace=True)
      traindf.drop(traindf[traindf['dropoff_latitude'] < -90].index, inplace=True)
```

### 4. Pickup and dropoff longitude should be ( -180 to 180 )

```
[22]: traindf.drop(traindf[traindf['pickup_longitude'] > 180].index, inplace=True)
      traindf.drop(traindf[traindf['pickup_longitude'] < -180].index, inplace=True)
      traindf.drop(traindf[traindf['dropoff_longitude'] > 180].index, inplace=True)
      traindf.drop(traindf[traindf['dropoff_longitude'] < -180].index, inplace=True)
```

```
[23]: traindf.shape
```

```
[23]: (14690, 12)
```

```
[24]: traindf.isna().sum()
```

```
[24]: fare_amount      0
      pickup_longitude  0
      pickup_latitude  0
      dropoff_longitude 0
      dropoff_latitude 0
      passenger_count   0
      Hour              0
      minute            0
      date              0
      day               0
      month             0
      year              0
      dtype: int64
```

#### 0.4.1 Calculate distance using Haversin formulas

```
[25]: from math import *

      def haversine(a):
          lon1=a[0]
          lat1=a[1]
          lon2=a[2]
          lat2=a[3]
          """
          Calculate the great circle distance between two points
          on the earth (specified in decimal degrees)
          """
          # convert decimal degrees to radians
          lon1, lat1, lon2, lat2 = map(radians, [lon1, lat1, lon2, lat2])
```

```

# haversine formula
dlon = lon2 - lon1
dlat = lat2 - lat1
a = sin(dlat/2)**2 + cos(lat1) * cos(lat2) * sin(dlon/2)**2
c = 2 * asin(sqrt(a))
# Radius of earth in kilometers is 6371
km = 6371* c
return km

```

```

[26]: traindf['distance'] =_
      ↪traindf[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude']].
      ↪apply(haversine,axis=1)

```

```

C:\Users\mukul\AppData\Local\Temp\ipykernel_20088\3890309117.py:4:
FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a
future version, integer keys will always be treated as labels (consistent with
DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
lon1=a[0]
C:\Users\mukul\AppData\Local\Temp\ipykernel_20088\3890309117.py:5:
FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a
future version, integer keys will always be treated as labels (consistent with
DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
lat1=a[1]
C:\Users\mukul\AppData\Local\Temp\ipykernel_20088\3890309117.py:6:
FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a
future version, integer keys will always be treated as labels (consistent with
DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
lon2=a[2]
C:\Users\mukul\AppData\Local\Temp\ipykernel_20088\3890309117.py:7:
FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a
future version, integer keys will always be treated as labels (consistent with
DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
lat2=a[3]

```

```

[27]: traindf.
      ↪drop(['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude'],_
      ↪axis=1, inplace=True)

```

```

[28]: traindf['Hour'] = traindf['Hour'].astype('int64')
traindf['minute'] = traindf['minute'].astype('int64')
traindf['date'] = traindf['date'].astype('int64')
traindf['day'] = traindf['day'].astype('int64')
traindf['month'] = traindf['month'].astype('int64')
traindf['year'] = traindf['year'].astype('int64')

```

```

[29]: traindf.head()

```



```
[29]:   fare_amount  passenger_count  Hour  minute  date  day  month  year  \
0         4.5           1.0      17      26    15   0     6    2009
1        16.9           1.0      16      52     5   1     1    2010
2         5.7           2.0       0      35    18   3     8    2011
3         7.7           1.0       4      30    21   5     4    2012
4         5.3           1.0       7      51     9   1     3    2010

      distance
0  1.030764
1  8.450134
2  1.389525
3  2.799270
4  1.999157
```

Distance should be positive and less than 130 Km.

```
[30]: traindf['distance'].sort_values(ascending=True).head(10)
```

```
[30]: 1542      0.0
      8135      0.0
      8130      0.0
      8123      0.0
      8109      0.0
      13446     0.0
      1397      0.0
      8068      0.0
      8063      0.0
      1419      0.0
      Name: distance, dtype: float64
```

```
[31]: traindf.drop(traindf[traindf['distance'] >= 130].index, inplace=True)
      traindf.drop(traindf[traindf['distance'] <= 0 ].index, inplace=True)
```

```
[32]: #traindf.drop(traindf[traindf['distance'].isna() ].index, inplace=True)
```

```
[33]: traindf.head()
```

```
[33]:   fare_amount  passenger_count  Hour  minute  date  day  month  year  \
0         4.5           1.0      17      26    15   0     6    2009
1        16.9           1.0      16      52     5   1     1    2010
2         5.7           2.0       0      35    18   3     8    2011
3         7.7           1.0       4      30    21   5     4    2012
4         5.3           1.0       7      51     9   1     3    2010


      distance
0  1.030764
1  8.450134
2  1.389525
```

```
3  2.799270
4  1.999157
```

### 0.4.2 EDA


```
[34]: plt.figure(figsize=(10,5))
      sns.boxplot(y=traindf['passenger_count'])
```

```
[34]: <Axes: ylabel='passenger_count'>
```




```
[35]: plt.figure(figsize=(10,5))
      sns.boxplot(y=traindf['fare_amount'])
```

```
[35]: <Axes: ylabel='fare_amount'>
```



```
[36]: sns.boxplot(y=traindf['distance'])
```

```
[36]: <Axes: ylabel='distance'>
```



```
[37]: traindf.describe()
```

```
[37]:
```

	fare_amount	passenger_count	Hour	minute	\
count	14244.000000	14244.000000	14244.000000	14244.000000	
mean	11.310145	1.320604	13.490312	29.746069	
std	9.434057	0.675145	6.514905	17.288493	
min	-3.000000	0.000000	0.000000	0.000000	
25%	6.000000	1.000000	9.000000	15.000000	
50%	8.500000	1.000000	14.000000	30.000000	
75%	12.500000	1.000000	19.000000	45.000000	
max	108.000000	4.000000	23.000000	59.000000	

	date	day	month	year	distance
count	14244.000000	14244.000000	14244.000000	14244.000000	14244.000000
mean	15.668211	3.030820	6.264462	2011.739890	3.451184
std	8.685281	1.971422	3.444883	1.870919	4.674488
min	1.000000	0.000000	1.000000	2009.000000	0.000111
25%	8.000000	1.000000	3.000000	2010.000000	1.277104
50%	16.000000	3.000000	6.000000	2012.000000	2.196772
75%	23.000000	5.000000	9.000000	2013.000000	3.935514
max	31.000000	6.000000	12.000000	2015.000000	129.950482

```
[38]: traindf['passenger_count'] = traindf['passenger_count'].astype('int64')
```

```
[39]: traindf
```

```
[39]:
```

	fare_amount	passenger_count	Hour	minute	date	day	month	year	\
0	4.5	1	17	26	15	0	6	2009	
1	16.9	1	16	52	5	1	1	2010	
2	5.7	2	0	35	18	3	8	2011	
3	7.7	1	4	30	21	5	4	2012	
4	5.3	1	7	51	9	1	3	2010	
...	...	...	...	...	...	...	...	...	
16062	6.5	1	7	41	12	4	12	2014	
16063	16.1	2	7	58	13	0	7	2009	
16064	8.5	1	11	19	11	2	11	2009	
16065	8.1	1	23	53	11	1	5	2010	
16066	8.5	2	6	24	14	2	12	2011	

	distance
0	1.030764
1	8.450134
2	1.389525
3	2.799270
4	1.999157
...	...
16062	0.850044
16063	7.867638
16064	1.469105

```
[14244 rows x 9 columns]
```

```
[41]: temp =traindf['passenger_count'].count()
      x=0
      z=[]
      for x in range(temp):
          g=traindf['passenger_count'].iloc[x]
          j=traindf['fare_amount'].iloc[x]
          z.append(passenger(g,j))
      traindf['Changed_fare']=z
```

```
[42]:
```

	fare_amount	passenger_count	Hour	minute	date	day	month	year	\
0	4.5	1	17	26	15	0	6	2009	
1	16.9	1	16	52	5	1	1	2010	
2	5.7	2	0	35	18	3	8	2011	
3	7.7	1	4	30	21	5	4	2012	
4	5.3	1	7	51	9	1	3	2010	
...	...	...	...	...	...	...	...	...	
16062	6.5	1	7	41	12	4	12	2014	
16063	16.1	2	7	58	13	0	7	2009	
16064	8.5	1	11	19	11	2	11	2009	
16065	8.1	1	23	53	11	1	5	2010	
16066	8.5	2	6	24	14	2	12	2011	

	distance	Changed_fare
0	1.030764	90.0

1	8.450134	338.0
2	1.389525	124.0
3	2.799270	154.0
4	1.999157	106.0
...	...	...
16062	0.850044	130.0
16063	7.867638	332.0
16064	1.469105	170.0
16065	2.590036	162.0
16066	3.898113	180.0

[14244 rows x 10 columns]

```
[43]: traindf.isna().sum()
```

```
[43]: fare_amount      0
passenger_count      0
Hour                 0
minute              0
date                0
day                 0
month               0
year                0
distance             0
Changed_fare        373
dtype: int64
```

### 0.4.3 Train Test Split

```
[44]: traindf=traindf.dropna()
```

```
[45]: traindf.head()
```

```
[45]:   fare_amount  passenger_count  Hour  minute  date  day  month  year  \
0         4.5             1      17      26    15   0     6   2009
1        16.9             1      16      52     5   1     1   2010
2         5.7             2       0      35    18   3     8   2011
3         7.7             1       4      30    21   5     4   2012
4         5.3             1       7      51     9   1     3   2010

   distance  Changed_fare
0  1.030764          90.0
1  8.450134         338.0
2  1.389525         124.0
3  2.799270         154.0
4  1.999157         106.0
```

```
[46]: X=traindf.drop(['fare_amount', 'Changed_fare'],axis=1)
      y=traindf['Changed_fare'].astype('int64')
```

```
[47]: X
```

```
[47]:
```

	passenger_count	Hour	minute	date	day	month	year	distance
0	1	17	26	15	0	6	2009	1.030764
1	1	16	52	5	1	1	2010	8.450134
2	2	0	35	18	3	8	2011	1.389525
3	1	4	30	21	5	4	2012	2.799270
4	1	7	51	9	1	3	2010	1.999157
...	...	...	...	...	...	...	...	...
16062	1	7	41	12	4	12	2014	0.850044
16063	2	7	58	13	0	7	2009	7.867638
16064	1	11	19	11	2	11	2009	1.469105
16065	1	23	53	11	1	5	2010	2.590036
16066	2	6	24	14	2	12	2011	3.898113

[13871 rows x 8 columns]

```
[48]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,
      ↪random_state = 42)
```

#### 0.4.4 Model selection

```
[49]: from sklearn.ensemble import RandomForestRegressor

      model = RandomForestRegressor(n_estimators=150,
      ↪criterion='absolute_error',max_depth=11)
```

```
[50]: model.fit(X_train, y_train)
```

```
[50]: RandomForestRegressor(criterion='absolute_error', max_depth=11,
      n_estimators=150)
```

```
[51]: predicts = model.predict(X_test)
```

```
[52]: model.score(X_train, y_train)
```

```
[52]: 0.903482640567547
```

#### 0.4.5 Model Evaluation

```
[53]: from sklearn.metrics import r2_score

      score=r2_score(y_test,predicts)
      score
```

```
[53]: 0.8294424351455236
```

```
[54]: from sklearn.metrics import mean_squared_error

RMSE = np.sqrt(mean_squared_error(y_test, predicts))
RMSE
```

```
[54]: 77.85002005591717
```

```
[55]: import pickle
filename = 'finalized_model.sav'
pickle.dump(model, open(filename, 'wb'))
```

```
[56]: loaded_model = pickle.load(open(filename, 'rb'))
result = loaded_model.score(X_test, y_test)
print(result)
```

```
0.8294424351455236
```

```
[57]: pred=loaded_model.predict(np.array([2,17,26,15,0,11,2022,1.030764])).
      ↪reshape(1,-1))
```

```
c:\Users\mukul\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\base.py:493: UserWarning: X does not have valid feature names,
but RandomForestRegressor was fitted with feature names
  warnings.warn(
```

```
[58]: pred
```

```
[58]: array([152.11333333])
```